

# Chapter 39 Files and Directories

So far we have covered **virtualization** of

- The CPU → process
- Memory (volatile) → the address space

Expand virtualization to **persistent storage (non-volatile memory)**

- keep stored content even after power loss
- Hard disk drive (HDD)
  - Block addressable
- Solid-state storage device (SSDs)
  - Block addressable
- Non-volatile memory (NV memory):
  - Optane (3D X-point)
  - byte addressable

- Root directory: /
- Separator: /
- Sub-directories
- The file system provides unified way to access
  - Files on disk
  - USB drives
  - CD-ROMs
- Absolute path names
  - /foo/bar/fizz.txt

## 31.1 Files and Directories

Two key abstractions for virtualization of storage

### 1. File

- Array of bytes
- Low-level name: inode identifier
- OS does not know the structure of the file
- OS responsibility:
  - Persistently store data and
  - Retrieve data again (without corruption)
- File name parts: bar.txt
  - Bar and txt (separate period .)
  - Just a convention; not enforced by the OS or file system

### 2. Directory

- Has structure
- Also has low-level name: inode number
- Contains entries (user-readable name, low-level name)
  - ("foo", 10)
- Each entry in a directory, either
  - A file
  - Other directory
- Placing directories in other directories
  - Can build arbitrary directory tree

## 39.3 Creating Files

To create a file, use the open system call:

```
int fd = open("foo", O_CREAT | O_WRONLY | O_TRUNC, S_IRUSR | S_IWUSR)
```

Second param:

- O\_CREAT - create file if does not exist
- O\_WRONLY - file can only be written to
- O\_TRUNC - if file already exists, truncate it to size zero (remove existing content)

Third param:

- S\_IRUSR - make file readable by owner (current user)
- S\_IWUSR - make file writable by owner

Return:

- fd - file descriptor
  - Integer
  - Private per process
  - Use fd to read or write the file (assuming you have permission)
  - The fd as just a handle or pointer that gives you power to perform certain operations
- fd: as pointer to an object of type file
- Call methods to access the file: read() and write()

Go has a similar function in the "os" package:

- os.OpenFile()

File descriptors managed by OS (per-process basis; kept in `proc` struct on Unix/xv6 kernel)

```
struct proc {  
    ...  
    struct file *ofile[NOFILE];    // open files  
    ...  
}
```

`NOFILE` = max # of files per-process

Each entry in array — just a pointer to a “file” struct.

Commands:

- `strace` (Linux)
- `dttruss` (macOS) — but must be enabled from recovery mode

Note on the `open()` call:

- `O_RDONLY` - file is opened as read-only by cat
- Returns success (positive value = `fd` = 3)

Q: Why 3?

Already has three files open: `stdin` = 0, `stdout` = 1, `stderr` = 2

You can see these fds in other calls to `read(3,...)` and `write(1, ...)`

`read(fd, content, size)` = size of content

`write(fd, content, size)` = size of content

`write(1, ...)` write content to `stdout`/screen

`Open()`: offset initialized to zero

`read(..., 100)`; - 100 bytes at a time

`read(..., 100)` returns 0 — done

`close(fd)`

OFT = Open File Table

Two entries (10, 11)

Two file descriptors (3, 4), but pointing to the same “file”

Current offset for each OFT entry is updated independently

Process uses `lseek()` to reposition current offset before reading (or writing)

- set current offset to 200
- Read next 50 bytes
  - Updates the current offset to 250

cat tries to read more from the file:

`read(3, “...”, size) = 0`

So this returns 0, meaning that it has read the entire file

cat calls

`close(3)`

### 39.5 Reading and Writing, But Not Sequentially

Now we want to read or write to a specific offset within a file.

`off_t lseek(int fd, off_t offset, int whence);`

- `fd` - file descriptor
- Offset — position/location within the file
- Whence - how seek is performed
  - `SEEK_SET` - offset is set to offset in bytes
  - `SEEK_CUR` - offset is set to current location + offset bytes
  - `SEEK_END` - offset is set to size of file + offset bytes

xv6:

```
struct file {  
    int ref;        // reference count  
    char readable;  // opened as readable  
    char writable;  // opened as writable  
    struct inode *ip; // points to underlying file  
    uint off;       // current offset  
}
```

### 39.7 Writing Immediately with `fsync()`

For performance reasons:

- buffer writes in memory
- At later time (e.g. after 5 seconds), issue write ops to storage device
- Appears to calling application that write are quick

Issue: data may be lost

- rare case, e.g. due to machine crash after `write()` call, but before the write to disk takes place

This guarantee is called “eventual” durability and is “acceptable” for some systems.

However, many systems require strong durability guarantees.

- e.g. database management system (DBMS)

To support such applications

- `fsync(int fd)`
- Force all dirty (not yet written) data to be stored on disk
- Returns when all writes are complete (can take a bit of time)
- Application can safely move on knowing the data has been persisted

Note:

- may also need to `fsync()` the surrounding directory that contains the file foo
- If foo newly created

- Ensure that file foo is durable part of the directory

### 39.8 Renaming Files

```
% strace mv foo bar
```

```
Calls: rename("foo", "bar") = 0
```

Guarantee of rename(old, new):

- atomic with respect to system crashes
  - File will either be named "old" or "new"
  - No in-between state can arise

Ex. File editor (e.g. emacs); editing file foo.txt

- inserting a new line of text or code in the middle of the file
- To guarantee the new (saved) file has original contents plus the new line, do this:

```
int fd = open("foo.txt.tmp", O_WRONLY | O_CREAT | O_TRUNC, S_IRUSR | S_IWUSR);
write(fd, buffer, size);      // write out new version of file
fsync(fd);
close(fd);
rename("foo.txt.tmp", "foo.txt");
```

Editor code:

- write new version of the file under a temporary name (foo.txt.tmp)
- Force it to disk with fsync()
  - Application is now certain that the new file metadata and contents are on disk
- Rename temporary file to original file's name: foo.txt
- Atomically swapping new file into place; (deleting old version) — atomic file update

Empty directories has two entries:

- "." (dot) - refers to itself
- ".." (Dot-dot) - refers to the parent directory

### 39.9 Getting Information About Files

File system keeps lots of info about each file it is storing:

- such data about files is called metadata
- Get metadata for a file, use
  - stat()
  - fstat() system calls

The stat struct contains key info:

- file size
- inode#
- Ownership info
- When the file was access/modified

Linux: % stat file

macOS: % stat -x file

### 39.10 Removing Files

```
% strace rm foo
```

```
unlink("foo") = 0
```

```
% rm -rf /
```

### 39.11 Making Directories

```
% strace mkdir foo
```

```
mkdir("foo", 0777) = 0
```

When created, the directory is considered empty.

### 39.17 Making and Mounting a File System

Q: How to assemble a full directory tree from many underlying file systems?

- First make file systems: mkfs
- Then mounting them to make their contents accessible: mount

Goal: To make a file system accessible within a *uniform file system tree*.

Mount:

- Take existing directory as target *mount point*.
- "Paste" (or map) new file system onto directory tree at that point.

Example:

- ext3 file system
- One device partition: /dev/sda1
  - Content of device file system: meling, lampport
- Want to mount this file system at mount point: /home/users

```
% mkdir -p /home/users
```

```
% mount -t ext3 /dev/sda1 /home/users
```

```
% ls /home/users/
```

```
meling lampport
```

```
% ls /usr/local/
```

Mount command:

- apfs - Apple file system (optimized for SSDs)
- Devfs - device file system
- On linux: (check on Unix machines...)
  - ext3: hard disk file system
  - **proc**: file system for accessing information about processes
  - tmpfs: file system for temporary files
  - nfs: network file system